

Async Dart and agent-assisted coding

Futures, Streams, and coding with an agent in the loop

Dinu-Ștefan Rusu

FILS, Universitatea Politehnica București · Autumn 2026

What you should leave with



Await without blocking

Futures, `async/await`, and `try/catch/finally` around every network or database call.



Escape heavy work

Recognize CPU-bound code that freezes the UI, and move it off the main thread with `compute()`.



Streams in the UI

`FutureBuilder` and `StreamBuilder` turn asynchronous values into widgets that rebuild on their own.



Work with an agent

Prompt it with real context, then review its output like any other pull request.

PART 1 · ASYNC DART

One thread, many things waiting

Futures, await, and error handling

One thread runs your whole UI

Every build, every layout pass and every gesture handler runs on the same thread as your Dart code.

A function that takes 400 ms to return blocks all three for 400 ms. The screen stops responding.

async and await exist so a function can wait for something without occupying that thread while it waits.

The event loop and microtask queue that make this precise belong to Mobile and Embedded Computing (MEC), Lecture 2.

Blocking is not a Flutter bug. It is what one thread does when nothing yields control back to it.

A Future is one value, later

```
Future<String> fetchName() {  
    return Future.delayed(  
        const Duration(seconds: 1),  
        () => 'Dinu',  
    );  
}  
  
void main() {  
    fetchName().then(print);  
    print('This runs first.');}
```

Future<T> is a promise.

The value is not there yet, but the type of the eventual value is known now.

`then` runs a callback when the value arrives. `await` is sugar that reads like synchronous code: next slide.

async and await: suspending without blocking

```
Future<User> loadUser(String id) async {  
  final r = await dio.get('/users/$id');  
  return User.fromJson(r.data);  
}
```

```
Future<void> refresh(String id) async {  
  final user = await loadUser(id);  
  setState(() => _user = user);  
}
```

await reads top to bottom.

The compiler rewrites it into callbacks; you write it like synchronous code.

Each await is a suspension point: the function hands control back to the event loop and resumes when the Future completes.

Sequential awaits, or all at once

```
// Sequential: one after another
final user = await loadUser(id);
final posts = await loadPosts(id);

// Concurrent: both start now
final results = await Future.wait([
  loadUser(id),
  loadPosts(id),
]);
```

Sequential waits add up.

Two 200 ms calls in sequence cost about 400 ms.

`Future.wait` starts every `Future` immediately and completes when all of them do, or fails as soon as one does.

What happens when an awaited call fails

A failure inside async code does not throw at the call site. It completes the Future with an error instead.

await rethrows that error at the point of the await, so try, catch and finally behave exactly as they do in synchronous code.

Without a try around it, the error escapes as an unhandled async error. Users rarely report it; the screen just seems to do nothing.

Every await is a place your code can fail. Wrap it, or plan for the silence when you do not.

Handling errors around await

```
Future<void> refresh(String id) async {  
  try {  
    _user = await loadUser(id);  
  } on DioException catch (e) {  
    _error = 'Network error';  
  } catch (e, stack) {  
    _error = 'Unexpected error';  
    Crashlytics.recordError(e, stack);  
  } finally {  
    if (mounted) {  
      setState(() => _loading = false);  
    }  
  }  
}
```

One specific catch, one general catch.

Handle the failure you expect by type, then anything else.

Check `mounted` before calling `setState` after an `await`. The widget may already be gone.

PART 2 · STREAMS AND HEAVIER WORK

When one value is not enough

Streams, FutureBuilder, StreamBuilder, and compute()

Streams: many values, over time

```
// A Future is one value, later.  
final name = await fetchName();  
  
// A Stream is many values, over  
// time.  
Stream<int> ticker() async* {  
  for (var i = 0; ; i++) {  
    await Future.delayed(  
      const Duration(seconds: 1),  
    );  
    yield i;  
  }  
}
```

await for reads a Stream.

Or subscribe with `.listen` for a callback style.

Streams are how WebSockets and sensor feeds reach your UI. Cancel every subscription you create by hand, in `dispose()`.

Where these Futures and Streams come from

HTTP calls through dio or http return a Future. Lecture 7, networking.

Future.delayed and Timer model waiting without a real request.

sqlite and Drift queries return a Future; watching a table returns a Stream. Lecture 9, local storage.

Firestore snapshots and WebSocket messages arrive as Streams. Lecture 8 and Lecture 11.

FutureBuilder: showing a Future in the tree

```
FutureBuilder<User>(  
  future: _userFuture,  
  builder: (context, snapshot) {  
    if (snapshot.connectionState == ConnectionState.waiting) {  
      return const CircularProgressIndicator();  
    }  
    if (snapshot.hasError) return Text('Error: ${snapshot.error}');  
    return UserCard(user: snapshot.data!);  
  },  
)
```

`connectionState` tells you where the Future is: waiting, then done. Check `hasError` before reading `data`.

Do not rebuild the Future on every build

Recreated every build

```
class ProfilePage extends
  StatelessWidget {
  Widget build(context) {
    return FutureBuilder(
      future: fetchUser(),
      builder: (c, s) => Body(s),
    );
  }
}
```

Stored once, in initState

```
class _ProfilePageState
  extends State<Page> {
  late final Future<User>
    _future = fetchUser();
  Widget build(context) =>
    FutureBuilder(
      future: _future,
      builder: (c, s) => Body(s),
    );
}
```

Store the Future once. A widget that calls `fetchUser()` inside `build()` refetches on every rebuild.

StreamBuilder: showing a Stream in the tree

```
StreamBuilder<int>(  
  stream: ticker(),  
  initialData: 0,  
  builder: (context, snapshot) {  
    if (snapshot.hasError) {  
      return Text('Error: ${snapshot.error}');  
    }  
    return Text('Tick: ${snapshot.data}');  
  },  
)
```

StreamBuilder subscribes when the widget builds and cancels when it is removed. You never call `listen` yourself. `initialData` covers the gap before the first event.

Putting it together: load once, then stay live

Load what already exists with a Future: chat history, a cached list, the first page of results. Wrap that Future in a FutureBuilder so the screen shows a spinner, then the initial data. Subscribe to what changes with a Stream: new messages, a live price, a sensor reading. Nest a StreamBuilder inside, or beside, the FutureBuilder so new events update the same screen.

Most real screens need both. A Future for what is already true, a Stream for what keeps changing.

FutureBuilder vs StreamBuilder, at a glance

ASPECT

FUTUREBUILDER

STREAMBUILDER

Input	one Future	a Stream
Rebuilds	once, when it completes	on every new event
Cancels for you	nothing to cancel	yes, when the widget is removed
Snapshot carries	connectionState, data, error	connectionState, data, error
Typical source	an HTTP call, a database read	a WebSocket, a live query

Async is not parallel

```
// This function is async. It still freezes the UI.  
Future<int> sumTo(int n) async {  
    var total = 0;  
    for (var i = 0; i < n; i++) {  
        total += i;    // pure CPU work  
    }  
    return total;    // loop never yielded  
}
```

async only helps when something else is doing the waiting: a socket, a disk, a timer.

Rule of thumb: waiting on I/O calls for `await`; burning CPU for longer than a frame calls for `compute()`.

compute(): the one-call escape for heavy work

```
// Must be a top-level or static function.  
List<Photo> parsePhotos(String body) {  
    final list = jsonDecode(body) as List;  
    return list.map(Photo.fromJson).toList();  
}  
  
Future<List<Photo>> fetchPhotos() async {  
    final res = await http.get(url);           // I/O  
    return compute(parsePhotos, res.body);     // CPU  
}
```

compute() spawns a short-lived isolate, runs one function and sends the result back. Arguments and results must be sendable: primitives, Strings, Lists and Maps.

Rule of thumb: `await` or `compute()`

Waiting on a network call, a disk read, a timer or a database calls for `await`. Something else is doing the work; your thread stays free to keep drawing frames.

Burning CPU for longer than a frame calls for `compute()`. Nothing is waiting; the loop only gets a turn back when the function returns.

Spawning an isolate by hand with `Isolate.spawn`, the full event loop, and how Dart's isolates compare to a Go goroutine are MEC Lecture 2.

`compute()` is the escape you will reach for in this course. What is happening underneath it is next semester's material.

PART 3 · AGENT-ASSISTED CODING

Coding with an agent in the loop

The 2026 tool landscape, and the review discipline it requires

The 2026 tool landscape



Agentic CLIs

Claude Code, Codex CLI, Gemini CLI: they run in your repository, edit files and run your tests.



IDE assistants

Cursor, GitHub Copilot, JetBrains AI: completion plus in-editor chat.



Cloud agents

Long-running tasks against a checkout; what comes back is a pull request you review.



Review bots

Automated tools comment on your pull request before a human ever opens it.



Tool access (MCP)

Agents reach your files, shell, database and APIs through a tool protocol.



One layer down

All of them are a shell over a handful of frontier models; the shell is what differs.

MCP: how an agent reaches your tools

MCP, the Model Context Protocol, is how an agent calls real tools instead of only generating text.

A typical agentic CLI can read and edit files, run your test suite, query a database, or call an internal API, all through the same protocol.

That access is not free: a misread instruction can touch files or data you never intended to expose.

Tool access raises the stakes of a wrong prompt. The rest of this section is about keeping that risk under control.

What agents are good at

Boilerplate: scaffolding, data classes, test fixtures, glue code.

Explaining unfamiliar code, stack traces and build errors.

Mechanical refactors that touch many files at once.

The test you were quietly going to skip.

A first draft of a task you are unsure how to start.

Where agents cost you

Plausible but wrong APIs: Flutter's surface moves fast.

Security holes: hardcoded keys, unvalidated input, permissive database rules.

Code that ignores conventions you never told it about.

Silent drift away from the package versions you are actually on.

The error paths and edge cases nobody asked for.

An agent produces confident output whether or not it is correct. Everything it writes still lands in the commit log under your name.

Where this bites, specifically in Flutter

An agent trained on older material may suggest `google_generative_ai` when `firebase_ai` is the current path. Lecture 12.

It may reach for `setState` inside a class that is not a `State`, or skip `const` where it would help.

It may hardcode an API key directly into a Dart file, which ships inside the compiled app. Lecture 7 and Lecture 10.

It rarely writes the empty-list, no-network or expired-token branch unless you ask for it by name.

None of this means do not use it. It means read the diff with the same suspicion you would give a stranger's pull request.

Prompting: give it context, not wishes

Name the files, packages and versions the change has to fit into.

Break multi-part work into numbered steps: one concern per prompt.

State your conventions explicitly. It cannot guess your codebase.

Ask for a diff or a test, not just a description.

Then read the diff yourself.

```
// Weak
add pagination to the list

// Better
In feed_page.dart, convert ListView
to ListView.builder. Paginate: 20
items per page, fetch next page at
200px from bottom, show a spinner.
```

Conventions: Dio for HTTP, no
setState outside State, const
where possible. Add a widget test.

Review it like any other pull request



Read every line

If you cannot explain it in review, you cannot ship it, or defend it at the practical test.



Check the security surface

Hardcoded keys, unvalidated input, permissive database rules.



Test the edges

Empty list, no network, expired token: the paths nobody thought to ask for.



Check the API is current

Models learn from old docs.
Verify the package and method still exist.



Keep the slices small

One concern per branch and pull request; generated code is only reviewable in slices.



You own it

“The agent wrote it” is not a defense in code review, and not one here either.

Using agents in this course's labs: the rules

RULE

WHAT IT MEANS IN PRACTICE

Allowed	Use any agent or assistant you like, for every lab and the semester app.
Must be understood	You must be able to explain any generated line if asked, live, during the lab.
Must be reviewed	Read the diff before you commit it, exactly like a teammate's pull request.
Cited in the PR	Name the tool and what you asked it to do, in the pull request description.

This is not a trust exercise: the same review discipline every professional team applies.

What citing an agent in a pull request looks like

```
## What changed
Added pagination to the feed screen.

## Agent use
Claude Code drafted the ListView.builder
conversion and the loading-row widget
test from a written prompt. I reviewed
every line, fixed one wrong Dio option,
and added the empty-state case myself.
```

This is what *must be reviewed* and *cited in the PR* look like together: specific about what the agent did, specific about what you checked yourself.

Three things to remember

1. `await` is not a thread. It suspends one function and frees the event loop until the `Future` completes.
2. A CPU-bound loop still blocks the UI. Reach for `compute()` the moment work outlasts one frame.
3. An agent's output is confident, not necessarily correct. You review it, you understand it, you own it.

FURTHER READING

dart.dev/language/async · [FutureBuilder cookbook](#)

Questions?

dinu_stefan.rusu@upb.ro · Microsoft Teams: course channel