

Offline-first sync

An outbox, a drain loop, and one conflict rule

Dinu-Ștefan Rusu

FILS, Universitatea Politehnica București · Week 10

What this lab is for



Write once, sync later

Every note write commits to the notes table and an outbox row, in one transaction. No edit is ever lost to a dropped connection.

TIME

2 hours



Resolve one conflict, visibly

Two devices editing the same note converge on the server's version, and the device that lost the race shows a marker.

SUBMIT

Push to your team repo, branch `lab-5`, before next lab

Where your code goes

1. Branch the repository, add Drift, and generate the boilerplate:

```
git checkout -b lab-5
flutter pub add drift sqlite3_flutter_libs
flutter pub add -d drift_dev build_runner
dart run build_runner build \
  --delete-conflicting-outputs
```

2. Local code lives in `lib/lab5/`. Server code lives in `server/lab5/main.go`.

Regenerate after every schema change

Adding a column to Notes or Outbox needs a fresh build and a bumped schema version, or an existing install crashes on the next launch.

Schema and outbox, one transaction

1. Model `Notes`: `id`, `title`, `body`, `updatedAt`, `version`, `isDirty`, `deletedAt`.
2. Model `Outbox`: `id`, `entityId`, `op`, `payload`, `queuedAt`.
3. Write `saveNote` so both tables commit inside one `db.transaction`.
4. Make delete a write: set `deletedAt`, never issue a SQL DELETE.

DONE WHEN

- ☐ Editing a note updates `notes` and inserts one `outbox` row.
- ☐ Killing the app mid-save leaves either both rows, or neither.
- ☐ Deleting a note still leaves an `outbox` row for the delete.

A write is not durable until the outbox has it too. An `outbox` row inserted outside the transaction can vanish if the app dies between the two writes.

Two tables, one purpose

```
class Notes extends Table {  
  TextColumn get id => text();  
  TextColumn get title => text();  
  TextColumn get body => text();  
  DateTimeColumn get updatedAt =>  
    dateTime();  
  IntColumn get version =>  
    integer();  
  BoolColumn get isDirty =>  
    boolean();  
  DateTimeColumn get deletedAt =>  
    dateTime().nullable();  
}
```

```
class Outbox extends Table {  
  TextColumn get id => text();  
  TextColumn get entityId => text();  
  TextColumn get op => text();  
  TextColumn get payload => text();  
  DateTimeColumn get queuedAt =>  
    dateTime();  
}
```

Both writes, one transaction

```
Future<void> saveNote(Note n) => db.transaction(() async {  
    final row = n.copyWith(  
        updatedAt: DateTime.now(), isDirty: true);  
    await into(notes).insertOnConflictUpdate(row);  
    await into(outbox).insert(OutboxCompanion.insert(  
        id: uuid.v4(), entityId: n.id,  
        op: 'upsert', payload: jsonEncode(row.toJson()),  
    ));  
});
```

watch() on notes fires the moment this commits. The outbox row leaves later, but the edit is already on screen.

Drain the outbox, then pull the delta

1. Add `drainOutbox()`: read the outbox oldest first, POST each payload.
2. On success, delete the row and clear `isDirty` in one transaction.
3. On start, and when connectivity returns, call `pull()` with the last saved version.
4. Trigger both from a connectivity listener, never a timer.

DONE WHEN

- ☐ Airplane mode on, edit, airplane mode off: the row leaves the outbox within a few seconds.
- ☐ Restarting with a non-empty outbox drains it before the list refreshes.
- ☐ `pull()` never re-downloads a row it already has.

The drain loop

```
Future<void> drainOutbox() async {  
  final batch = await (select(outbox)  
    ..orderBy([(o) => OrderingTerm.asc(o.queuedAt)]))  
    .get();  
  for (final item in batch) {  
    final r = await http.post(Uri.parse('$base/notes'), body: item.payload);  
    if (r.statusCode != 200) break;    // stop, keep order  
    await db.transaction(() async {  
      await (delete(outbox)..where((o) => o.id.equals(item.id))).go();  
      await notesDao.clearDirty(item.entityId);  
    });  
  }  
}
```


Delta sync on start

```
Future<void> pull() async {  
    final since = await meta.lastVersion();  
    final uri = Uri.parse('$base/notes?since=$since');  
    final rows = jsonDecode((await http.get(uri)).body) as List;  
    var newest = since;  
    await db.transaction(() async {  
        for (final row in rows) {  
            await notesDao.upsertFromServer(row);  
            newest = max(newest, row['version'] as int);  
        }  
    });  
    await meta.setLastVersion(newest);  
}
```

The whole sync server, part 1

```
type Note struct{ ID, Title, Body, UpdatedAt, DeletedAt string; Version int }
var mu sync.Mutex; var version int
var notes = map[string]Note{}
func upsert(w http.ResponseWriter, r *http.Request) {
    var n Note
    json.NewDecoder(r.Body).Decode(&n)
    mu.Lock()
    version++
    n.Version = version
    notes[n.ID] = n
    mu.Unlock()
    json.NewEncoder(w).Encode(n)
}
```

The whole sync server, part 2

```
func list(w http.ResponseWriter, r *http.Request) {
    since, _ := strconv.Atoi(r.URL.Query().Get("since"))
    mu.Lock()
    defer mu.Unlock()
    var out []Note
    for _, n := range notes {
        if n.Version > since {
            out = append(out, n)
        }
    }
    json.NewEncoder(w).Encode(out)
}
```

main() wires POST and GET on /notes to these handlers, then listens on :8080.

One conflict rule, and a visible marker

1. On pull, if the local row is not dirty, upsert the server's row: no conflict.
2. If the local row is dirty and the incoming version is higher, the server wins: keep its value.
3. Mark the losing row `updatedElsewhere` instead of silently overwriting it.
4. Clear the marker once the user opens the note.

DONE WHEN

- ☐ Two devices editing the same note end with one value on both.
- ☐ The device that lost the race shows an "Updated elsewhere" badge.
- ☐ The badge clears once the note is opened.

The loser is not the older edit. It is the device that had not yet seen the version that won.

The rule: server version wins

```
Future<void> onIncoming(Map row) => db.transaction(() async {  
    final local = await notesDao.find(row['id'] as String);  
    if (local == null || !local.isDirty) {  
        return notesDao.upsertFromServer(row);    // no conflict  
    }  
    if ((row['version'] as int) > local.version) {  
        await notesDao.markUpdatedElsewhere(local.id);  
        return notesDao.upsertFromServer(row);    // we lost  
    }  
    // local.version >= incoming: we are still ahead, keep ours  
});
```

Both devices run this same function, so both reach the same answer without talking to each other.

The marker in the list

```
ListTile(  
  title: Text(note.title),  
  trailing: note.updatedElsewhere  
    ? const Chip(label: Text('Updated elsewhere'))  
    : null,  
  onTap: () => notesDao.clearFlag(note.id),  
)
```

The marker is the whole UI cost of this rule. No merge screen, no dialog: one boolean and one badge tell the user their edit did not survive.

What the grader will do

THE DEMO

- ☐ Airplane mode on, edit a note, airplane mode off: the server has the edit within a few seconds.
- ☐ Two devices, or two emulators, edit the same note. Both end up with the same value.
- ☐ The device that lost the race shows the "Updated elsewhere" marker.

IN THE REPOSITORY

The Drift schema, the sync engine, and `server/lab5/main.go`, with a README naming your conflict rule.

One commit per task. A screenshot of both devices showing the same value is the cheapest proof you can leave.

Four failures you will meet

Target of URI hasn't been generated: `app_db.g.dart`

`build_runner` has not run, or ran before the table file was saved. Re-run it, then restart analysis.

Connection refused, or the request hangs until it times out

The emulator's `localhost` is itself. Point the client at `10.0.2.2`, not `localhost`.

The outbox row count never drops to zero

An exception is swallowed silently, or the loop never runs because the connectivity check never fires. Log every iteration.

A device keeps re-downloading rows it already has

The checkpoint was saved before the transaction committed, or the comparison used `>=` instead of `>`.

Push before you leave.

Branch lab-5 · one commit per task · your conflict rule named in the README